

DATA_ENG300_HW1_Aircraft-1

January 26, 2026

1 Homework 1: Aircraft inventory analysis

In this homework, we will explore the aircraft inventory dataset from 2006 to 2023 (Source: Bureau of Transportation Statistics, www.bts.gov).

Some useful resources: - You may find explanation of the dataset attributes available in [BTS dataset attribute summary](#). - You may verify some of your analyses by comparing your results with the BTS analysis dashboard. The dashboard is reachable via the “Analysis” column in the attribute summary linked above.

General rules of thumb for homeworks: - Read the homework questions carefully. - Explain your choices. - Present your findings concisely. - Use tables, plots, and summary statistics to aid your presentation of findings. - If you have an idea in mind but could not implement (in code), present the idea thoroughly and how you would have implemented the code.

1.0.1 Tasks:

For all tasks below, create one or more functions for any systematic changes to the dataset. For instance, someone could run the sequence of functions to transform the dataset.

Specify the sequence of functions.

1. Investigate the missing data in this dataset. Specifically, for each of the following variables that have missing data, decide if any imputation is possible. Give your reasoning and code for how you decide to impute missing values.
 - Columns for investigation: CARRIER, CARRIER_NAME, MANUFACTURE_YEAR, NUMBER_OF_SEATS, CAPACITY_IN_POUNDS, and AIRLINE_ID.
 - For example, watch out for “North American Airlines” aircrafts. Are the CARRIER/UNIQUE_CARRIER column *really* missing?
2. Inspect the columns MANUFACTURER, MODEL, AIRCRAFT_STATUS, and OPERATING_STATUS. Decide, for each column, if transformation or standardization of data are required. Give your reasoning and code if you decide to transform the data.

Hints:

- For very messy data like manufacturer/model names, give your best attempt. It is okay to not catch them all.
- Use `value_counts()` to identify “big wins”.
- Break down into multiple steps, instead of having one line of code to do them all.

3. Remove data rows that still have missing values. Report the amount of remaining data you obtained.
4. Transformation and derivative variables
 - For the columns `NUMBER_OF_SEATS` and `CAPACITY_IN_POUNDS`, check the skewness in the variable and plot a histogram for each variable.
 - The Box-Cox transformation (`scipy.stats.boxcox`) is one possible way to transform variables into a “more-normal-like” variable. Apply the Box-Cox transformation for these two columns and save them as new columns, i.e. `XXXXXXXX_BOXCOX`.
 - Plot a histogram for each transformed variable.
 - Describe what you observe before and after transformation.
5. Feature engineering
 - Create a new column `SIZE` by the quartiles of `NUMBER_OF_SEATS`
 - below 25% percentile: `SMALL`
 - 25% - 50% percentile: `MEDIUM`
 - 50% - 75% percentile: `LARGE`
 - above 75% percentile: `XLARGE`
 - For each size group, provide and plot the proportions of aircrafts that are operating versus not (`OPERATING_STATUS`).
 - For each size group, provide and plot the proportions of aircrafts belonging to each aircraft status group (`AIRCRAFT_STATUS`).
 - Provide a written summary of your findings.

1.0.2 Dataset

You may read in the dataset with the following code. Otherwise, you can find the `dataset` available in the `dataset` directory. Note that `pandas.read_csv()` is capable of reading a zipped .csv file and no unzipping is required.

```
[137]: import pandas as pd
inventory = pd.read_csv('T_F41SCHEDULE_B43.csv')

/var/folders/hh/2dsrbb3d021_w0zs6jdygn480000gn/T/ipykernel_35512/2180530456.py:2
: DtypeWarning: Columns (11) have mixed types. Specify dtype option on import or
set low_memory=False.
    inventory = pd.read_csv('T_F41SCHEDULE_B43.csv')
```

```
[138]: inventory.info()

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 132313 entries, 0 to 132312
Data columns (total 17 columns):
 #   Column           Non-Null Count  Dtype  
 ---  --  
 0   YEAR            132313 non-null   int64  
 1   CARRIER          132264 non-null   object  
 2   CARRIER_NAME     132208 non-null   object  
 3   MANUFACTURE_YEAR 132310 non-null   float64 
 4   UNIQUE_CARRIER_NAME 132208 non-null   object
```

```

5   SERIAL_NUMBER      132313 non-null  object
6   TAIL_NUMBER        132313 non-null  object
7   AIRCRAFT_STATUS    132313 non-null  object
8   OPERATING_STATUS   132313 non-null  object
9   NUMBER_OF_SEATS    132306 non-null  float64
10  MANUFACTURER      132313 non-null  object
11  AIRCRAFT_TYPE     102380 non-null  object
12  MODEL              132302 non-null  object
13  CAPACITY_IN_POUNDS 132212 non-null  float64
14  ACQUISITION_DATE   131346 non-null  object
15  AIRLINE_ID         132208 non-null  float64
16  UNIQUE_CARRIER     132149 non-null  object
dtypes: float64(4), int64(1), object(12)
memory usage: 17.2+ MB

```

1.0.3 Task 1

- Investigate the missing data in this dataset. Specifically, for each of the following variables that have missing data, decide if any imputation is possible. Give your reasoning and code for how you decide to impute missing values.
 - Columns for investigation: CARRIER, CARRIER_NAME, MANUFACTURE_YEAR, NUMBER_OF_SEATS, CAPACITY_IN_POUNDS, and AIRLINE_ID.
 - For example, watch out for “North American Airlines” aircrafts. Are the CARRIER/UNIQUE_CARRIER column *really* missing?

[139]: # check if this is the right type of solution

```

def fix_carrier_name(dataframe):
    """
    For North American Airlines, fixes the CARRIER and the UNIQUE_CARRIER
    columns.

    Pandas has automatically converted 'NA' to null values.
    Replace these null values with 'NAA' for North American Airlines

    If UNIQUE_CARRIER_NAME is 'North American Airlines' and UNIQUE_CARRIER is
    null, then set UNIQUE_CARRIER to 'NAA'.

    If CARRIER_NAME is 'North American Airlines' and CARRIER is null, then set
    CARRIER to 'NAA'.
    """

    df = dataframe.copy()
    # query and change specific rows to fill in missing data
    df.loc[(df['UNIQUE_CARRIER_NAME'] == 'North American Airlines') &
           (df['UNIQUE_CARRIER'].isnull()), 'UNIQUE_CARRIER'] = 'NAA'
    df.loc[(df['UNIQUE_CARRIER_NAME'] == 'North American Airlines') &
           (df['UNIQUE_CARRIER'] == 'NA'), 'UNIQUE_CARRIER'] = 'NAA'

```

```

df.loc[(df['CARRIER_NAME'] == 'North American Airlines') & (df['CARRIER'].isnull()), 'CARRIER'] = 'NAA'
df.loc[(df['CARRIER_NAME'] == 'North American Airlines') & (df['CARRIER'] == 'NA'), 'CARRIER'] = 'NAA'

return df

```

```

[140]: def investigate(column):
    """
    Investigates a column of a dataframe.
    Prints out basic information.
    Checks if the column is numeric to see if missing data can be imputed.
    """

    # print basic information
    print(f"Column: {column.name}")
    print(f"Data type: {column.dtype}")
    print(f"Sample values: {column.unique()[:5]}")

    # why would we not be able to impute missing data? if non-numeric
    # or if we can fill in from other columns
    if column.dtype == 'object':
        print('Cannot use knn to impute missing data for non-numeric columns\n' +
              'with this setup.\n')
    return

    print(f"Number of missing values: {column.isnull().sum()}")


def impute_numeric_knn(df):
    """
    Given a dataframe, selects the numeric columns and imputes missing data.
    Uses KNN imputation with 5 neighbors.
    Plots the imputed data in individual bar charts for each column.
    Returns imputed dataframe
    """

    from sklearn.impute import KNNImputer

    numeric_columns = df.select_dtypes('number').columns

    imputer = KNNImputer(n_neighbors=5)
    df_numeric = df[numeric_columns]
    imputed_data = imputer.fit_transform(df_numeric)
    df_imputed = df.copy()
    df_imputed[numeric_columns] = imputed_data

```

```

# print number of imputed values for each column
for column in numeric_columns:
    num_imputed = df[column].isnull().sum()
    print(f"Column: {column}, Number of imputed values: {num_imputed}")

# plot original data on top of imputed data
import matplotlib.pyplot as plt
import seaborn as sns
for column in numeric_columns:

    # plots (but hard to tell if not a lot of imputed data)
    fig, ax = plt.subplots(1, 1)
    sns.histplot(df_imputed[column], stat='probability', alpha = 0.4, bins=30, color='red', label='knn imputation')
    sns.histplot(df[column], stat='probability', alpha = 0.6, bins = 30, color='blue', label='unimputed')
    ax.legend()
    plt.show()

return df_imputed

```

[141]: # choose from `CARRIER`, `CARRIER_NAME`, `MANUFACTURE_YEAR`, `NUMBER_OF_SEATS`,
↪ `CAPACITY_IN_POUNDS`, and `AIRLINE_ID`

```
investigate(inventory['AIRLINE_ID'])
```

Column: AIRLINE_ID
Data type: float64
Sample values: [20397. 20408. 19917. 20007. 19874.]
Number of missing values: 105

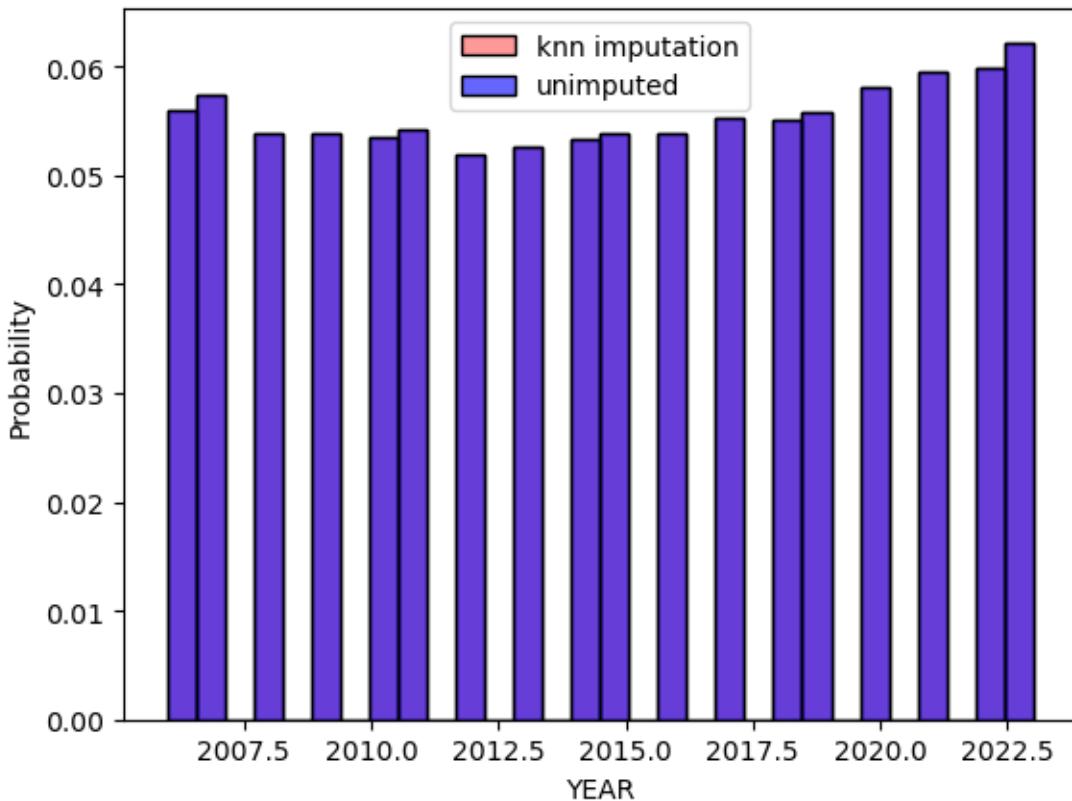
[142]: # print number of missing values in each column
print(inventory.select_dtypes('number').isnull().sum())

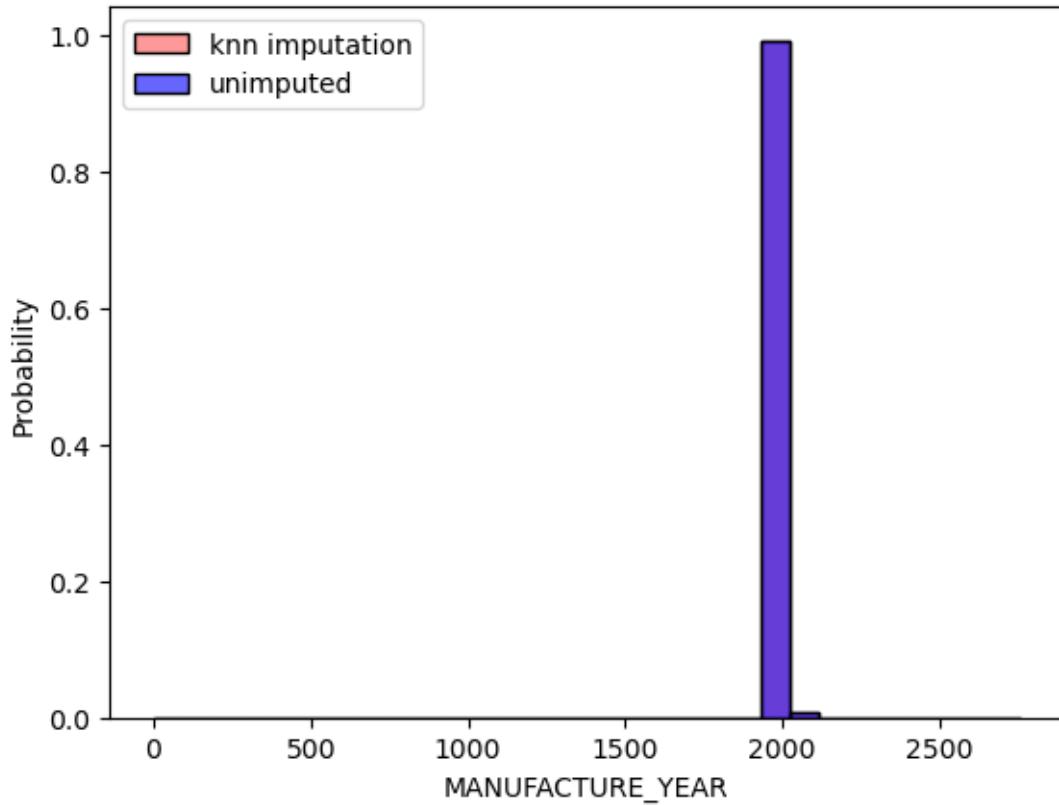
YEAR	0
MANUFACTURE_YEAR	3
NUMBER_OF_SEATS	7
CAPACITY_IN_POUNDS	101
AIRLINE_ID	105
dtype:	int64

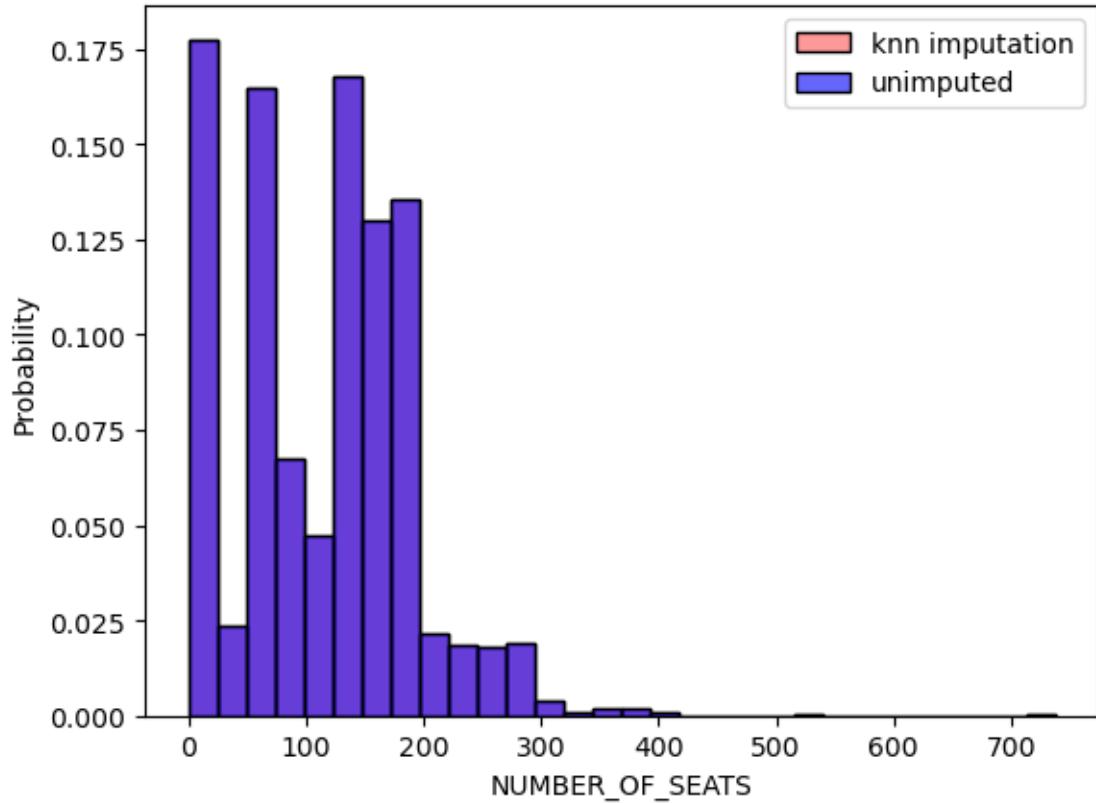
[143]: inventory_cleaned = fix_carrier_name(inventory)
inventory_imputed = impute_numeric_knn(inventory_cleaned)

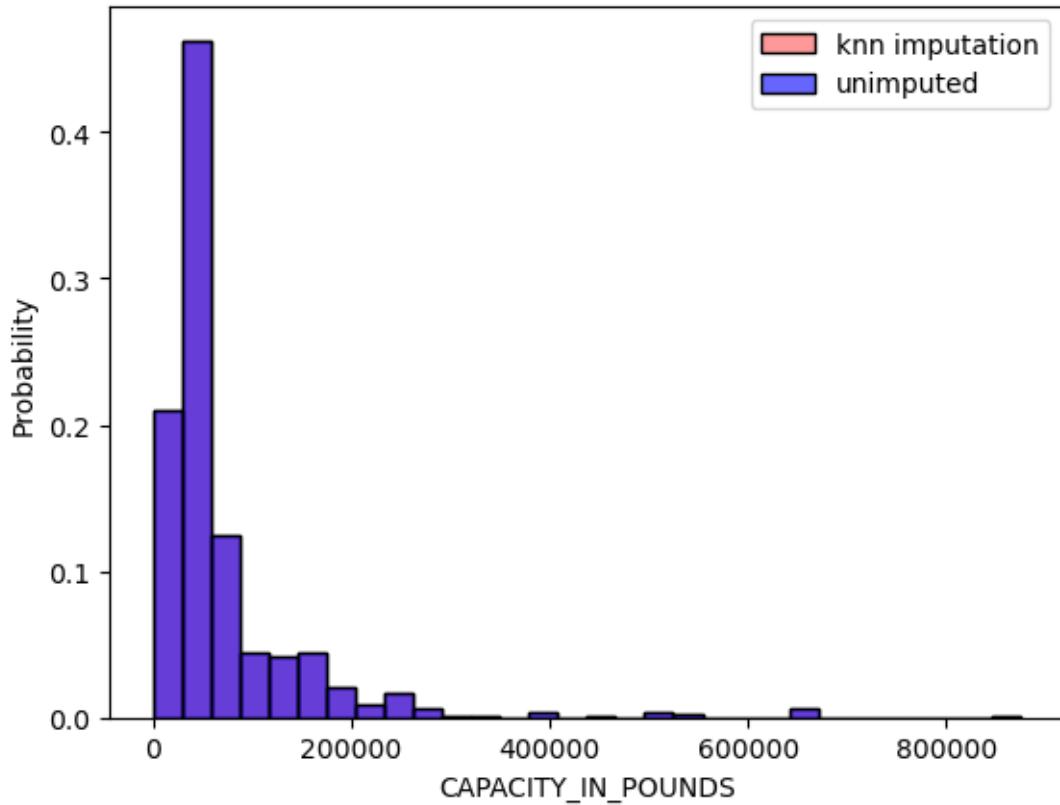
Column: YEAR, Number of imputed values: 0
Column: MANUFACTURE_YEAR, Number of imputed values: 3
Column: NUMBER_OF_SEATS, Number of imputed values: 7

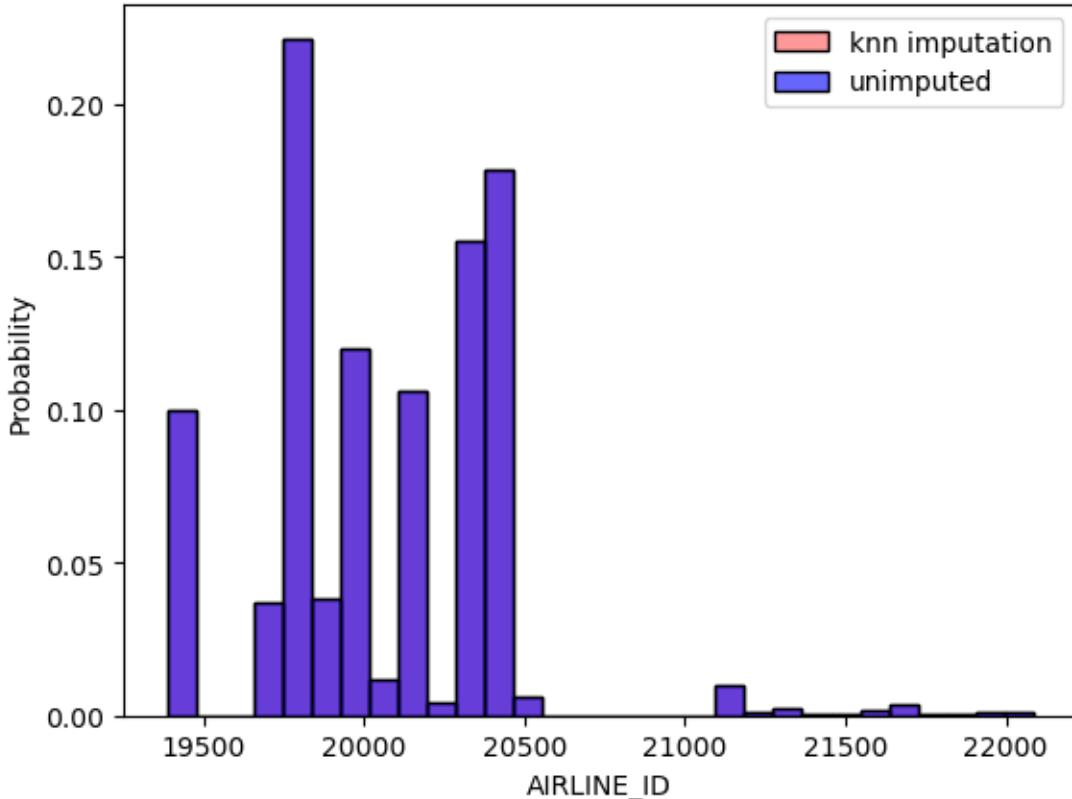
Column: CAPACITY_IN_POUNDS, Number of imputed values: 101
Column: AIRLINE_ID, Number of imputed values: 105











Response I decided to impute missing values in two ways: 1. For CARRIER and CARRIER_NAME, I saw that North American Airlines' code was 'NA'. When loading the dataset with pandas, it identified these as NaN. For columns where UNIQUE_CARRIER_NAME was 'North American Airlines' and UNIQUE_CARRIER was null, then I set UNIQUE_CARRIER to 'NAA'. For columns where CARRIER_NAME was 'North American Airlines' and CARRIER was null, then I set CARRIER to 'NAA'. I noticed there were also cases where it was not null because it included a space after ('NA'). I also transformed these to 'NAA'.

2. For the knn imputation, I decided to impute for all the numeric columns in the dataset. 216 values were imputed in total from the numeric columns. From the histograms, the imputed data is not very visible because it makes up a very small percentage of each column.

1.0.4 Task 2

2. Inspect the columns MANUFACTURER, MODEL, AIRCRAFT_STATUS, and OPERATING_STATUS. Decide, for each column, if transformation or standardization of data are required. Give your reasoning and code if you decide to transform the data.

Hints:

- For very messy data like manufacturer/model names, give your best attempt. It is okay to not catch them all.
- Use `value_counts()` to identify “big wins”.

- Break down into multiple steps, instead of having one line of code to do them all.

```
[144]: inventory_imputed['OPERATING_STATUS'].value_counts()
#inventory_imputed['MANUFACTURER'].value_counts()
#inventory_imputed['MODEL'].value_counts()
#inventory_imputed['AIRCRAFT_STATUS'].value_counts()
```

```
[144]: OPERATING_STATUS
Y      126577
N      5664
y       71
1        1
Name: count, dtype: int64
```

```
[145]: def uppercase(column):
    """
    transforms a column to uppercase
    ex. for aircraft status:
    O or o -> O
    """
    return column.str.upper()

def empty_to_na(column):
    """
    replaces empty strings with NaN
    """
    return column.mask(column.str.strip() == '', pd.NA)

def std_manufacturers(column):
    """
    standardizes manufacturer names to a simple form
    all entries containing 'BOEING' should be standardized to 'BOEING'
    all entries containing 'AIRBUS' should be standardized to 'AIRBUS'
    removes extra spaces
    """
    col = column.str.upper() # make uppercase
    col = col.apply(lambda x: 'BOEING' if 'BOEING' in x else x)
    col = col.apply(lambda x: 'AIRBUS' if 'AIRBUS' in x else x)
    return col.str.strip()

# standardize the columns
inventory_imputed['OPERATING_STATUS_STD'] = uppercase(inventory_imputed['OPERATING_STATUS'])
inventory_imputed['OPERATING_STATUS_STD'] = empty_to_na(inventory_imputed['OPERATING_STATUS_STD'])
```

```

inventory_imputed['AIRCRAFT_STATUS_STD'] = uppercase(inventory_imputed['AIRCRAFT_STATUS'])
inventory_imputed['MODEL_STD'] = uppercase(inventory_imputed['MODEL'])
inventory_imputed['MANUFACTURER_STD'] = std_manufacturers(inventory_imputed['MANUFACTURER'])

```

Response: From my initial inspection, I was able to identify several types of errors in several of the given columns.

For aircraft status and operating status, there are lowercase versions of several of the possible statuses. From the documentation for both columns, this is shown to be an error, because only the uppercase codes are assigned meanings. For these two columns, I standardized them by changing the entire column to uppercase. Operating status also had one entry that was a single space (' ') and I replaced this with NaN.

For model, I did not have enough prior knowledge to identify specific errors, but I transformed the column to uppercase to catch those cases.

For manufacturer, I noticed that there were many variations of the same manufacturer. For example, Boeing was written as BOEING, Boeing, THEBOEINGCO, etc. For this column, I transformed everything to uppercase, and then for some noticeable cases, I was able to standardize further. For Boeing and Airbus, if these strings were in the entry, I replaced the entry with these standardized versions. I also noticed that there were a couple entries that had excess spacing, so I removed that.

1.0.5 Task 3

3. Remove data rows that still have missing values. Report the amount of remaining data you obtained.

```
[146]: def remove_missing_data_rows(df):
    """
    removes rows with any missing data from the dataframe
    """
    clean_df = df.copy()
    print(f'Number of missing rows: {df.isnull().sum().sum()}') # print number of
    # missing rows per column
    print(f'\nOriginal shape: {df.shape}')
    clean_df = clean_df.dropna(axis=0, how='any') # drop whole rows if they
    # have any missing data
    print(f'\nNew shape: {clean_df.shape}')
    return clean_df
```

```
[147]: inventory_imputed = remove_missing_data_rows(inventory_imputed)
```

Number of missing rows:

YEAR	0
CARRIER	0
CARRIER_NAME	105
MANUFACTURE_YEAR	0

```
UNIQUE_CARRIER_NAME      105
SERIAL_NUMBER            0
TAIL_NUMBER              0
AIRCRAFT_STATUS          0
OPERATING_STATUS          0
NUMBER_OF_SEATS          0
MANUFACTURER             0
AIRCRAFT_TYPE            29933
MODEL                    11
CAPACITY_IN_POUNDS        0
ACQUISITION_DATE         967
AIRLINE_ID                0
UNIQUE_CARRIER           105
OPERATING_STATUS_STD      1
AIRCRAFT_STATUS_STD       0
MODEL_STD                 11
MANUFACTURER_STD          0
dtype: int64
```

Original shape: (132313, 21)

New shape: (101304, 21)

[148]: inventory_imputed.info()

```
<class 'pandas.core.frame.DataFrame'>
Index: 101304 entries, 29239 to 132312
Data columns (total 21 columns):
 #   Column           Non-Null Count  Dtype  
 ---  --  
 0   YEAR             101304 non-null   float64 
 1   CARRIER           101304 non-null   object  
 2   CARRIER_NAME      101304 non-null   object  
 3   MANUFACTURE_YEAR  101304 non-null   float64 
 4   UNIQUE_CARRIER_NAME 101304 non-null   object  
 5   SERIAL_NUMBER     101304 non-null   object  
 6   TAIL_NUMBER       101304 non-null   object  
 7   AIRCRAFT_STATUS   101304 non-null   object  
 8   OPERATING_STATUS  101304 non-null   object  
 9   NUMBER_OF_SEATS   101304 non-null   float64 
 10  MANUFACTURER     101304 non-null   object  
 11  AIRCRAFT_TYPE     101304 non-null   object  
 12  MODEL             101304 non-null   object  
 13  CAPACITY_IN_POUNDS 101304 non-null   float64 
 14  ACQUISITION_DATE  101304 non-null   object  
 15  AIRLINE_ID        101304 non-null   float64 
 16  UNIQUE_CARRIER    101304 non-null   object  
 17  OPERATING_STATUS_STD 101304 non-null   object
```

```

18 AIRCRAFT_STATUS_STD    101304 non-null  object
19 MODEL_STD                101304 non-null  object
20 MANUFACTURER_STD        101304 non-null  object
dtypes: float64(5), object(16)
memory usage: 17.0+ MB

```

Response The original dataframe had 132,313 entries. After removing all rows that contained any missing data, I was left with 101,305 entries. 29,933 of the deletions were due to missing data in the AIRCRAFT_TYPE column.

1.0.6 Task 4

4. Transformation and derivative variables

- For the columns NUMBER_OF_SEATS and CAPACITY_IN_POUNDS, check the skewness in the variable and plot a histogram for each variable.
- The Box-Cox transformation (`scipy.stats.boxcox`) is one possible way to transform variables into a “more-normal-like” variable. Apply the Box-Cox transformation for these two columns and save them as new columns, i.e. XXXXXXXXX_BOXCX.
- Plot a histogram for each transformed variable.
- Describe what you observe before and after transformation.

```
[149]: def skewtest_and_hist(column):
    """
    uses scipy skewtest (if p < 0.05, column is skewed)
    plots histogram of a column
    """
    from scipy.stats import skewtest
    import matplotlib.pyplot as plt

    # skew test
    _, p = skewtest(column)
    if p < 0.05:
        print(f'{column.name} is significantly skewed')
    else:
        print(f'{column.name} is not significantly skewed')

    print(f'p value: {p}')

    # plot histogram
    plt.figure(figsize=(4,3))
    plt.title(f'{column.name}')
    plt.xlabel(column.name)
    plt.ylabel('Frequency')
    plt.hist(column, bins=30, alpha=0.7)

def boxcox_transform(column):
    """
    checks if data is positive
    """
    pass
```

```

if negative, throws an error
if contains zeros, adds 1E-12 to all zeros
apply the boxcox transform to a single column
returns transformed column
'''

from scipy.stats import boxcox

# check for positivity
if (column < 0).any():
    print(f'{column.name} contains negative values. cannot apply box-cox')
elif (column == 0).any():
    # only add 1E-12 to zeros
    print(f'{column.name} contains zeros. adding 1E-12 to zeros before box-cox')
    column[column == 0] += 1E-12
else:
    print(f'{column.name} is positive. applying box-cox')

transformed_col, _ = boxcox(column)
return transformed_col

```

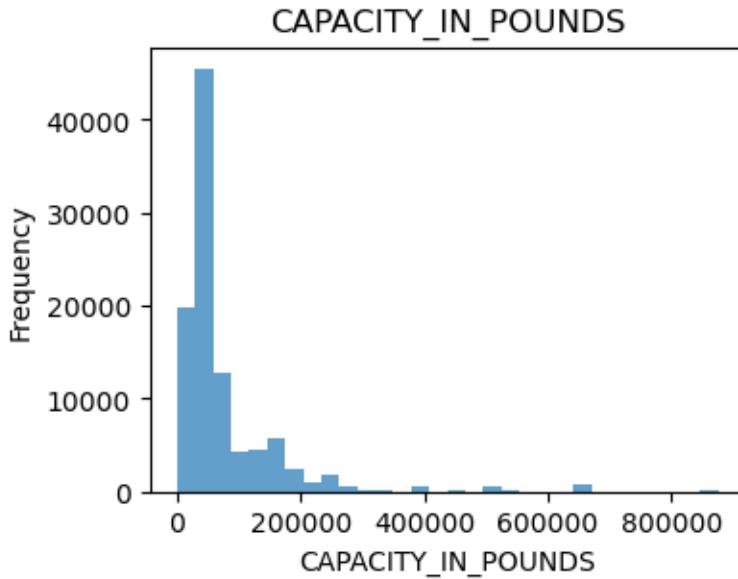
```
[150]: skewtest_and_hist(inventory_imputed['CAPACITY_IN_POUNDS'])
skewtest_and_hist(inventory_imputed['NUMBER_OF_SEATS'])
```

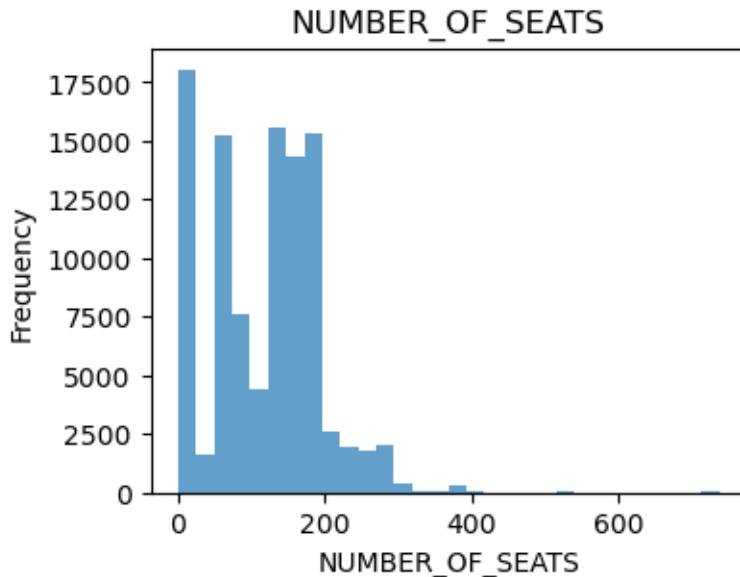
CAPACITY_IN_POUNDS is significantly skewed

p value: 0.0

NUMBER_OF_SEATS is significantly skewed

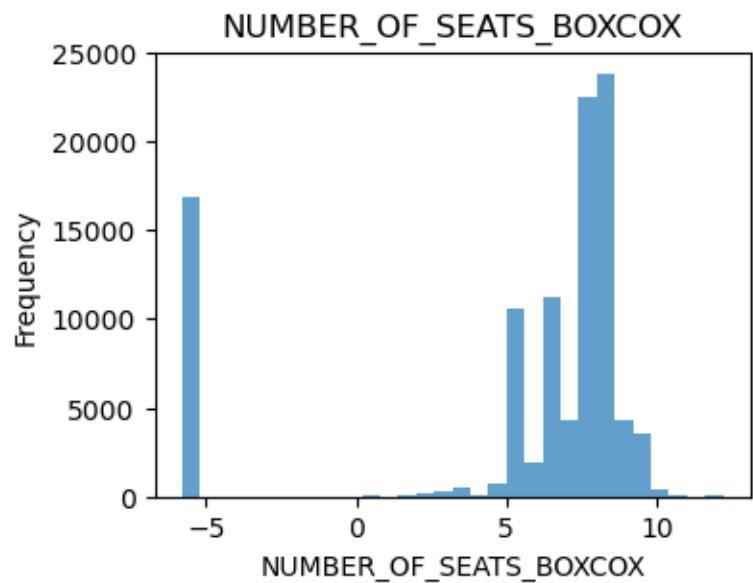
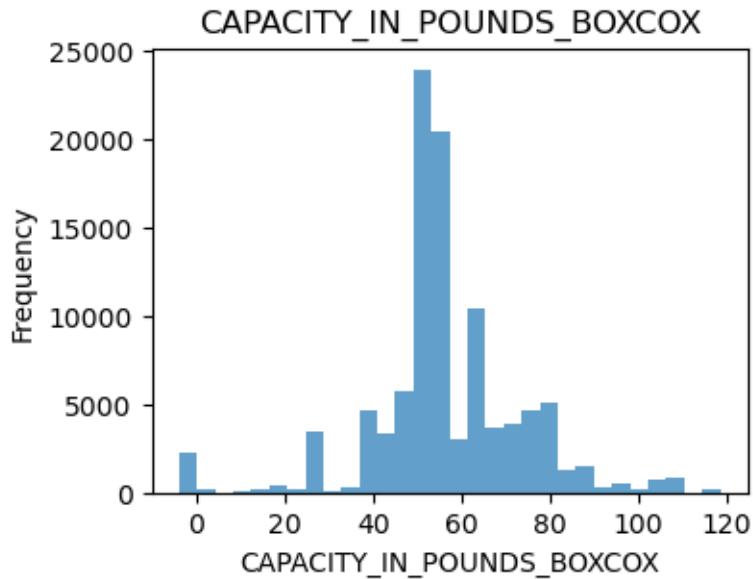
p value: 0.0





```
[151]: inventory_imputed['CAPACITY_IN_POUNDS_BOXCOX'] =  
    ~boxcox_transform(inventory_imputed['CAPACITY_IN_POUNDS'])  
skewtest_and_hist(inventory_imputed['CAPACITY_IN_POUNDS_BOXCOX'])  
  
inventory_imputed['NUMBER_OF_SEATS_BOXCOX'] =  
    ~boxcox_transform(inventory_imputed['NUMBER_OF_SEATS'])  
skewtest_and_hist(inventory_imputed['NUMBER_OF_SEATS_BOXCOX'])
```

CAPACITY_IN_POUNDS contains zeros. adding 1E-12 to zeros before box-cox
 CAPACITY_IN_POUNDS_BOXCOX is significantly skewed
 p value: 0.0
 NUMBER_OF_SEATS contains zeros. adding 1E-12 to zeros before box-cox
 NUMBER_OF_SEATS_BOXCOX is significantly skewed
 p value: 0.0



Response For both columns, I plotted a histogram and also performed a scipy skew test. Both of these analyses showed that the two columns were skewed. Because of this, I performed a boxcox transformation on the columns. This transformation requires that values be positive, but both columns had 0 values. For these specific values, I added 1E-12 so that the transformation could occur, without changing the value of the entry in a significant way.

From the post-box-cox histograms, the plots look much less skewed. This might be more useful if these data are being fed into a model or used for imputation of missing values. However, the values

also lose their meaning on their own unless they are converted back to the original values.

1.0.7 Task 5

5. Feature engineering

- Create a new column SIZE by the quartiles of NUMBER_OF_SEATS
 - below 25% percentile: SMALL
 - 25% - 50% percentile: MEDIUM
 - 50% - 75% percentile: LARGE
 - above 75% percentile: XLARGE
- For each size group, provide and plot the proportions of aircrafts that are operating versus not (OPERATING_STATUS).
- For each size group, provide and plot the proportions of aircrafts belonging to each aircraft status group (AIRCRAFT_STATUS).
- Provide a written summary of your findings.

```
[152]: def size_quartile_feature(df, old_feature):  
  
    print(f"Creating new feature 'SIZE' from '{old_feature}' based on quartiles.  
    ↵")  
  
    # create new feature  
    df['SIZE'] = None  
  
    # find quartiles  
    Q1 = df[old_feature].quantile(0.25)  
    Q2 = df[old_feature].quantile(0.50)  
    Q3 = df[old_feature].quantile(0.75)  
    print(f"Q1: {Q1}, Q2: {Q2}, Q3: {Q3}")  
  
    # assign new values based on quartiles  
    df.loc[df[old_feature] <= Q1, 'SIZE'] = 'SMALL'  
    df.loc[(df[old_feature] > Q1) & (df[old_feature] <= Q2), 'SIZE'] = 'MEDIUM'  
    df.loc[(df[old_feature] > Q2) & (df[old_feature] <= Q3), 'SIZE'] = 'LARGE'  
    df.loc[df[old_feature] > Q3, 'SIZE'] = 'EXTRA_LARGE'  
  
    # plot bar plot of new feature  
    import matplotlib.pyplot as plt  
    plt.figure(figsize=(6,4))  
    plt.title(f'Bar plot of "SIZE"')  
    plt.xlabel('SIZE')  
    plt.ylabel('Count')  
    df['SIZE'].value_counts().reindex(['SMALL', 'MEDIUM', 'LARGE',  
    ↵'EXTRA_LARGE']).plot(kind='bar')  
    plt.show()  
  
def plot_proportions(df, target_feature):
```

```

    ...
ex.
For each size group, provide and plot the proportions of aircrafts that are
→operating versus not (`OPERATING_STATUS`).
For each size group, provide and plot the proportions of aircrafts
→belonging to each aircraft status group (`AIRCRAFT_STATUS`).
    ...
import matplotlib.pyplot as plt

target_feature = df[target_feature]

size_order = ['SMALL', 'MEDIUM', 'LARGE', 'EXTRA_LARGE']
prop_table = (df.groupby('SIZE')[target_feature.name].
→value_counts(normalize=True).unstack(fill_value=0))
prop_table = prop_table.reindex(size_order)

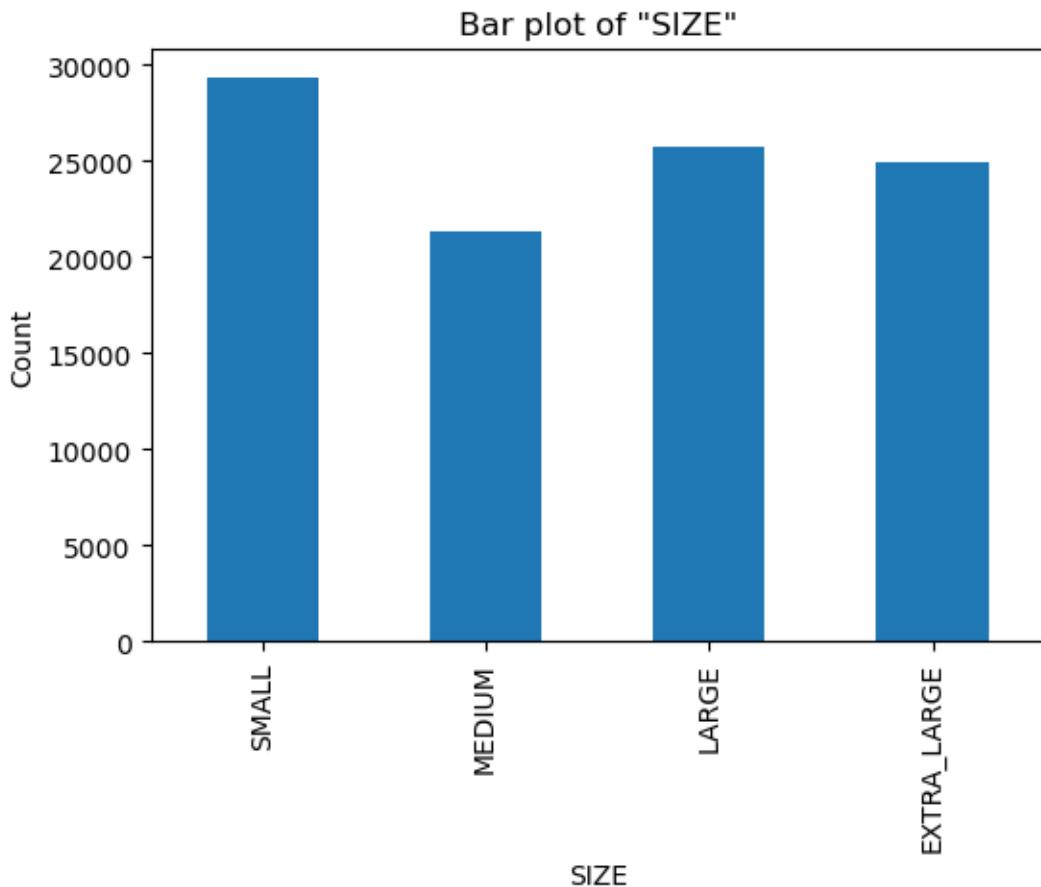
display(prop_table)

ax = prop_table.plot(kind='bar', figsize=(6,4), width=0.8)
ax.set_title(f'Proportion of {target_feature.name} by Size Group')
ax.set_ylabel('Proportion')
ax.set_xlabel('Size Group')
ax.legend(title=target_feature.name)
plt.show()

```

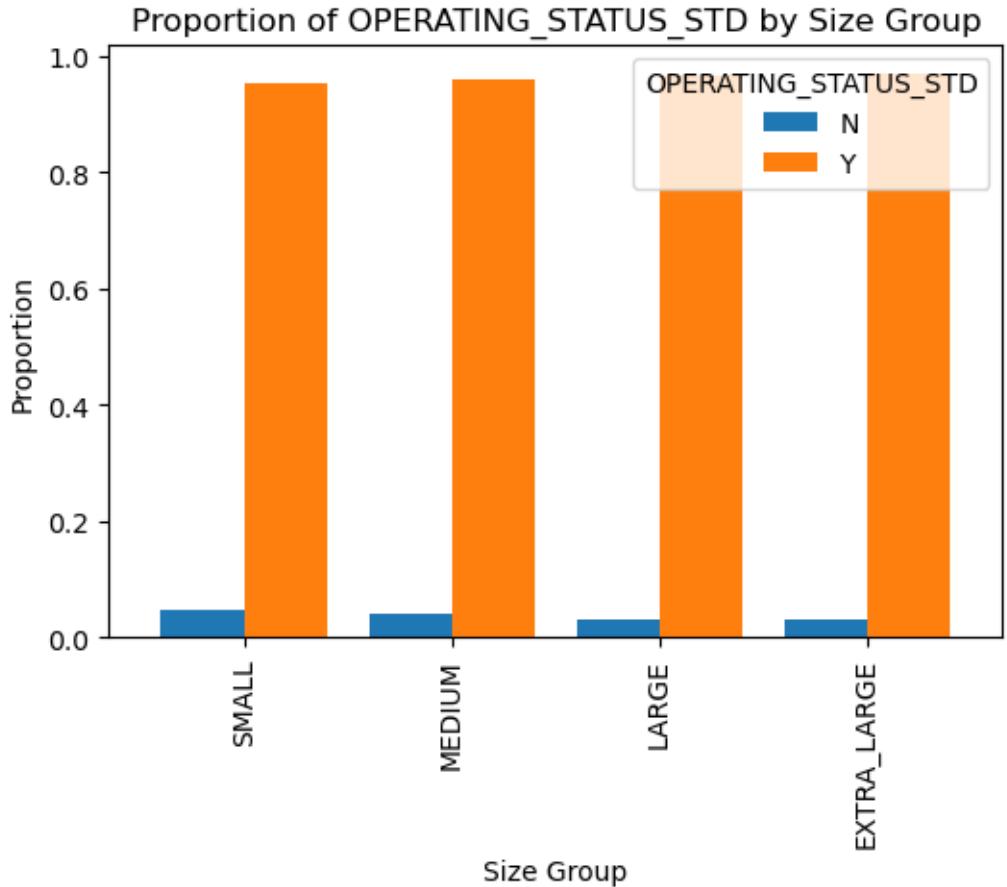
[153]: size_quartile_feature(inventory_imputed, 'NUMBER_OF_SEATS')

Creating new feature 'SIZE' from 'NUMBER_OF_SEATS' based on quartiles.
Q1: 50.0, Q2: 136.0, Q3: 169.0

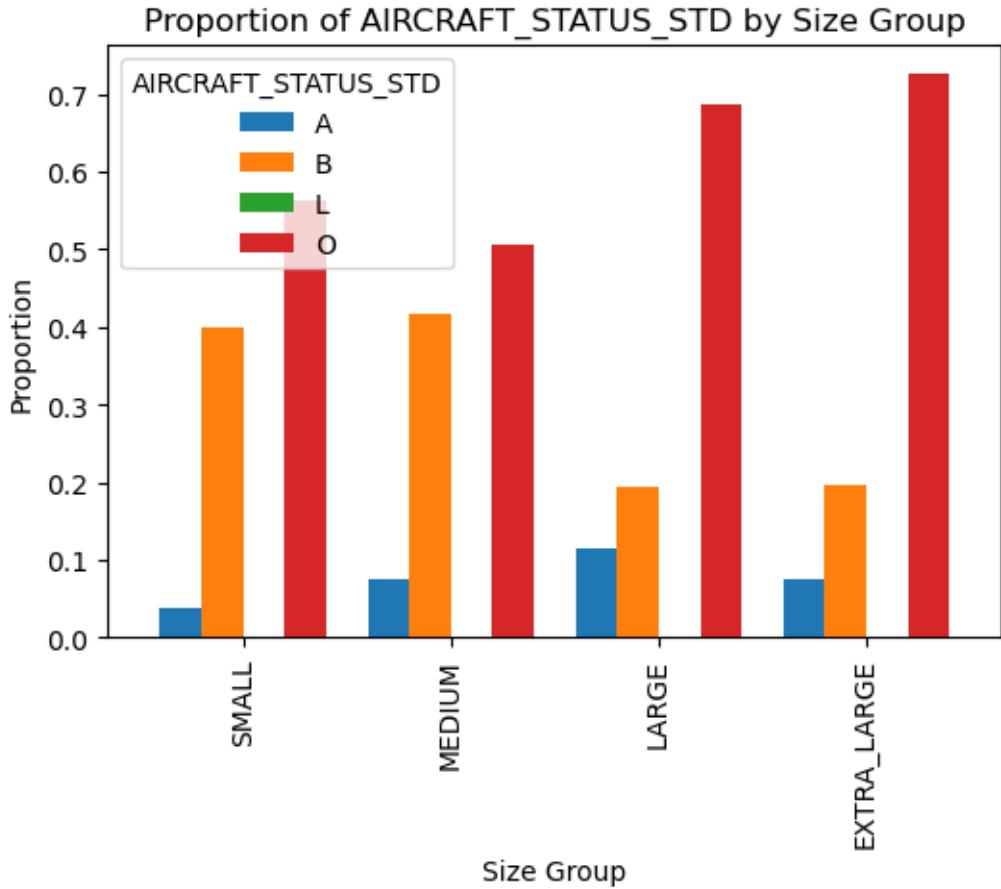


```
[154]: for feature in ['OPERATING_STATUS_STD', 'AIRCRAFT_STATUS_STD']:
    plot_proportions(inventory_imputed, feature)
```

	N	Y
OPERATING_STATUS_STD		
SIZE		
SMALL	0.046695	0.953305
MEDIUM	0.040401	0.959599
LARGE	0.032507	0.967493
EXTRA_LARGE	0.031157	0.968843



AIRCRAFT_STATUS_STD	A	B	L	O
SIZE				
SMALL	0.039327	0.398868	0.000000	0.561805
MEDIUM	0.076401	0.417115	0.001592	0.504892
LARGE	0.115545	0.195040	0.001869	0.687546
EXTRA_LARGE	0.075547	0.197169	0.001604	0.725680



Response From the plot of Operating Status by Size, it is visible that the proportion of ‘Y’ vs ‘N’ across size groups is relatively constant. However, from the proportion of Aircraft Status by Size, it seems like although the relative sizes are similar across groups, there are some differences. Primarily, the medium and small groups have a much larger proportion of ‘B’ status aircraft. The extra large and large groups have a smaller proportion of ‘B,’ and so they have larger proportions of ‘O.’ Because there are so few ‘L’ status aircraft, these are not visible on the plot.

2 Submission guidelines

Your submission should be contained in a `homework_1` folder of your Github repository, and it should include - a `readme.md` file including how to run the code and what your expected outputs are (if the code is run), - your source code, and - a `.pdf` or `.html` file containing your written answers.

3 Generative AI disclosure

Syllabus policy:

Required disclosure: each submission must include an AI Usage note stating: (1) tool(s) used, (2) the key prompt(s), and (3) what you changed and how you verified the results. If none, write: “AI Usage: None.”

I used ChatGPT to help generate the plotting function for 5 because I was struggling to show the grouping correctly.

Prompt: “Help complete this plot:

```
plot_proportions(dataframe, target_feature): """ ex. For each size group, provide and
plot the proportions of aircrafts that are operating versus not (OPERATING_STATUS).
For each size group, provide and plot the proportions of aircrafts belonging to each
aircraft status group (AIRCRAFT_STATUS). """
import matplotlib.pyplot as plt
target_feature = df[target_feature]
plt.figure(figsize=(6,4))
plt.title(f'{target_feature.name} by size group')
grouped = inventory_imputed.groupby('SIZE')[target_feature]
plt.bar(grouped.size(), grouped.size().values)
plt.xlabel('Size Group')
plt.ylabel('Proportion')
plt.show()
```

I was able to visually confirm that the plots looked as I was intending. I also chose to display the prop_table to check that the data made sense with what I would have expected. I also adjusted some small visual settings and added the reindexing.